

# Comparative analysis of ANN & LeNet-5 on MNIST dataset.

• Name - Rishi Sharma  
• RollNo - 17TC0135

In this notebook we will discuss the implementation of traditional **Artificial Neural Network(ANN)** & **LeNet-5**(a convolutional based Neural Network). Then we will compare their performance based on the **Accuracy**, **Precision** & **Recall** of two models.

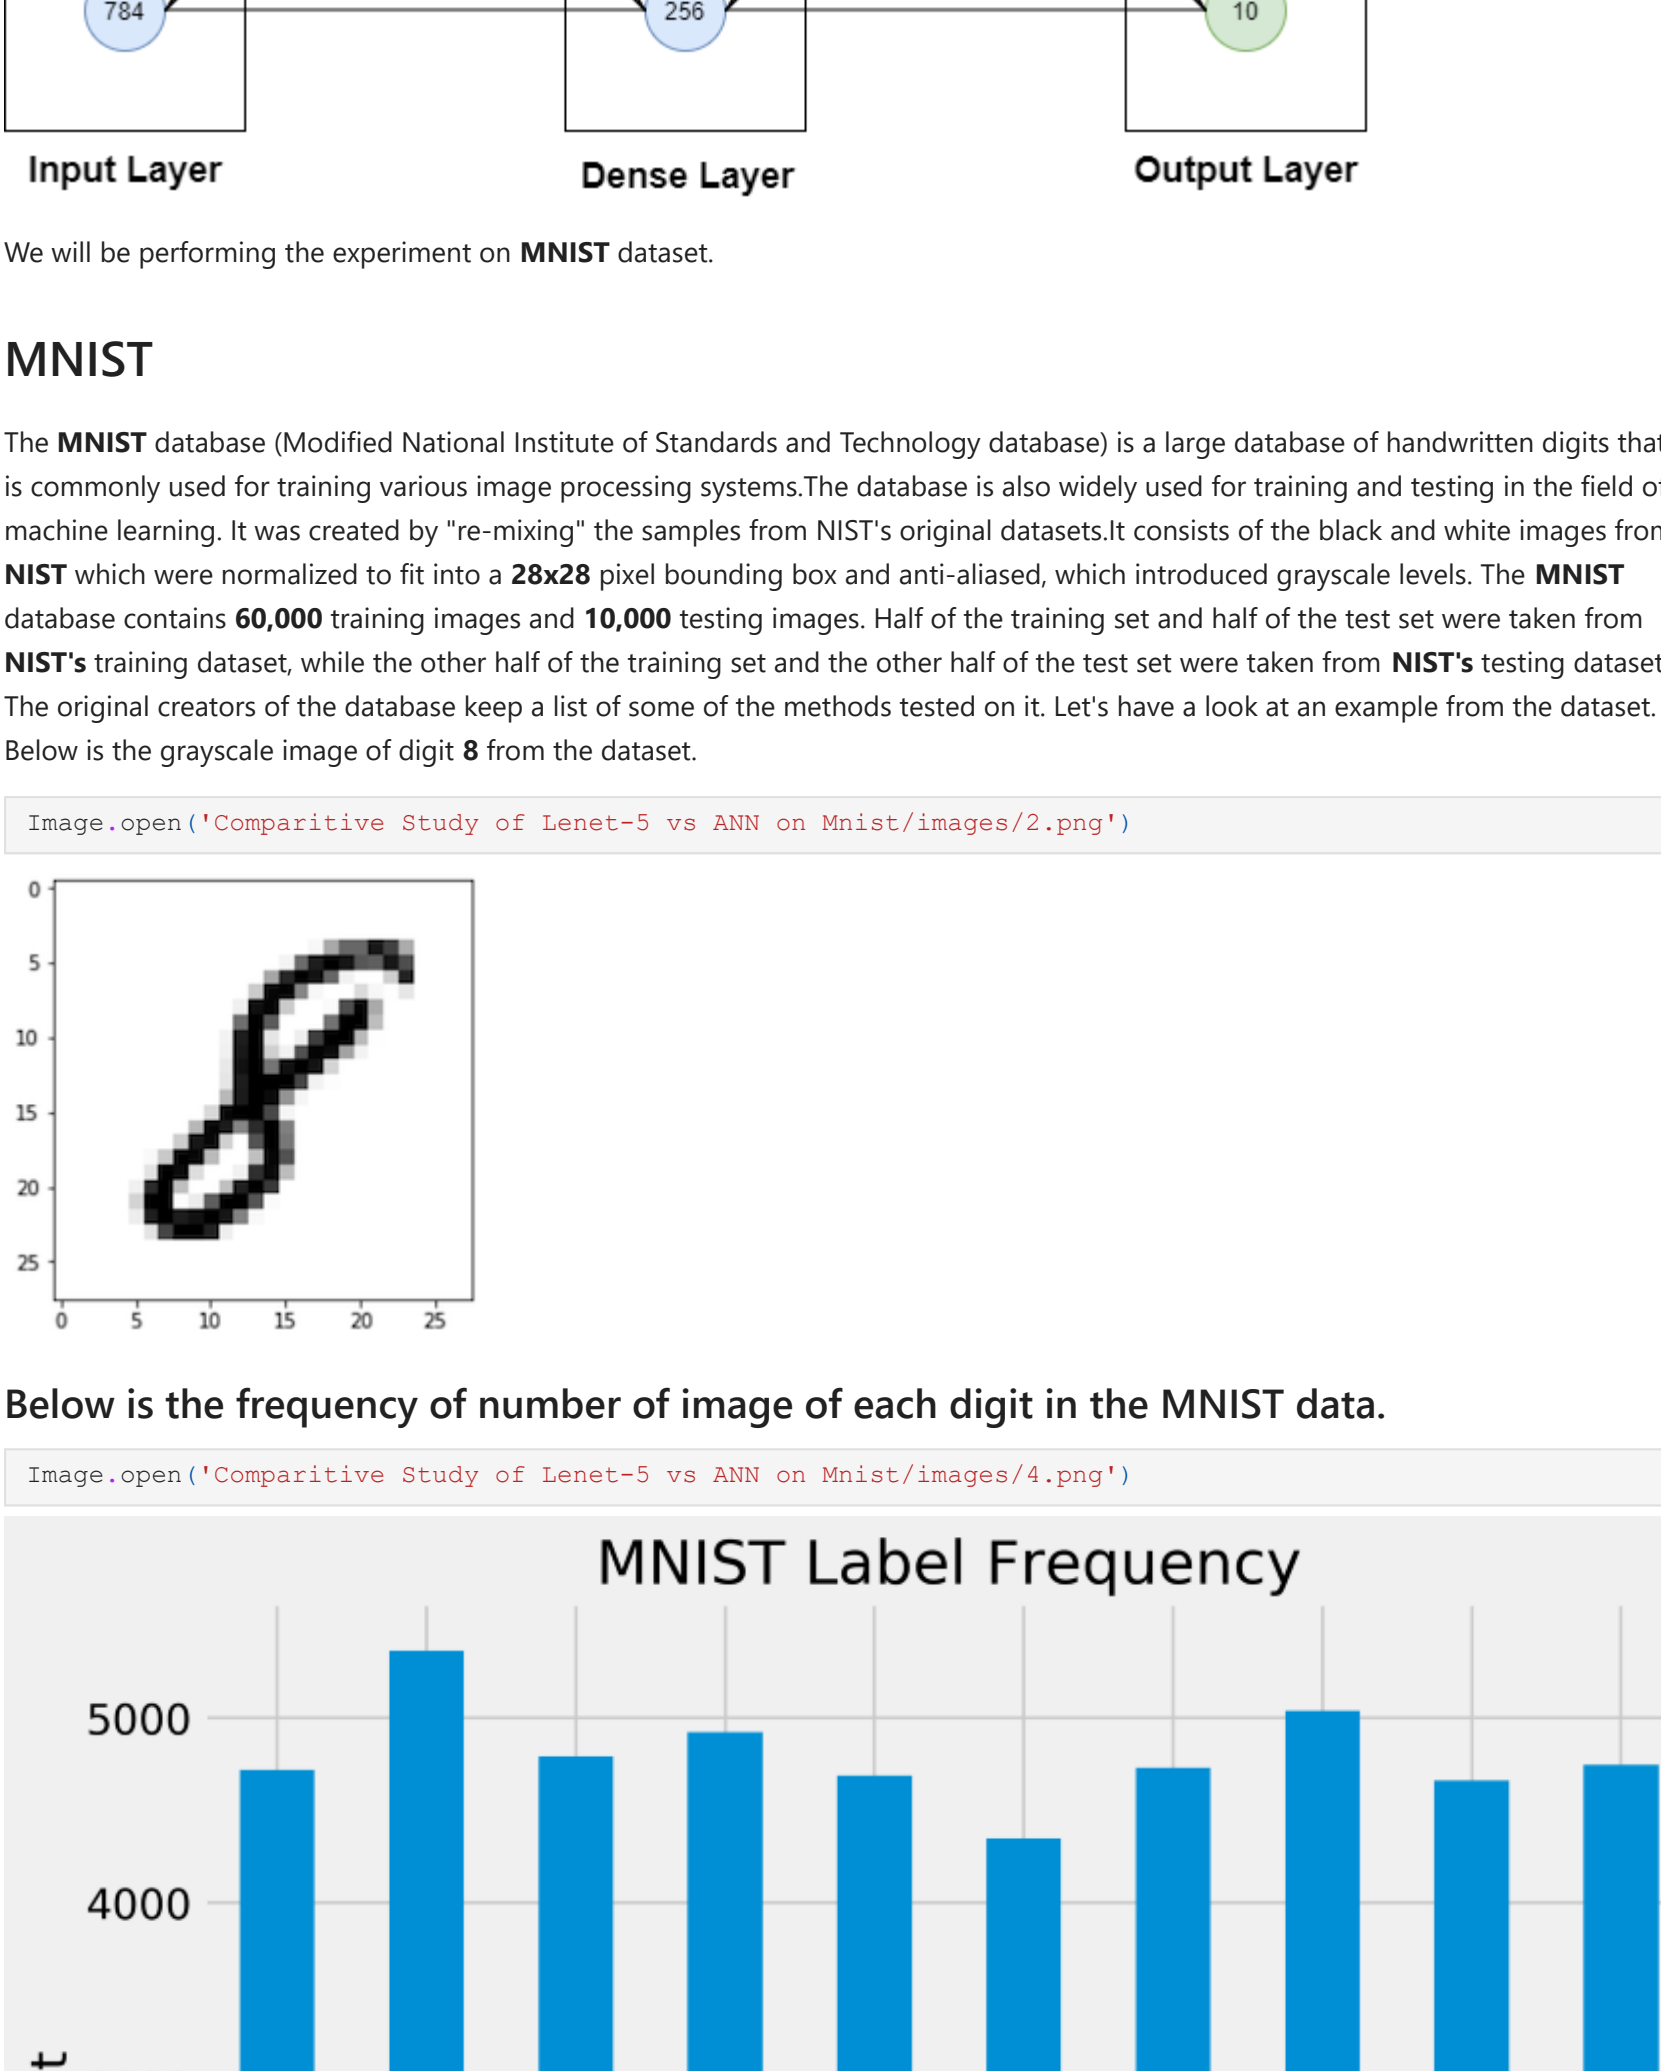
## Artificial Neural Networks(ANN)

An **Artificial Neural Network (ANN)** is a computer system inspired by biological neural networks for creating artificial brains based on the collection of connected units called artificial neurons. It is designed to analyse and process information as humans. Artificial Neural Network has self-learning capabilities to produce better results as more data is available.

An **Artificial Neural Network (ANN)** is composed of four principal objects:

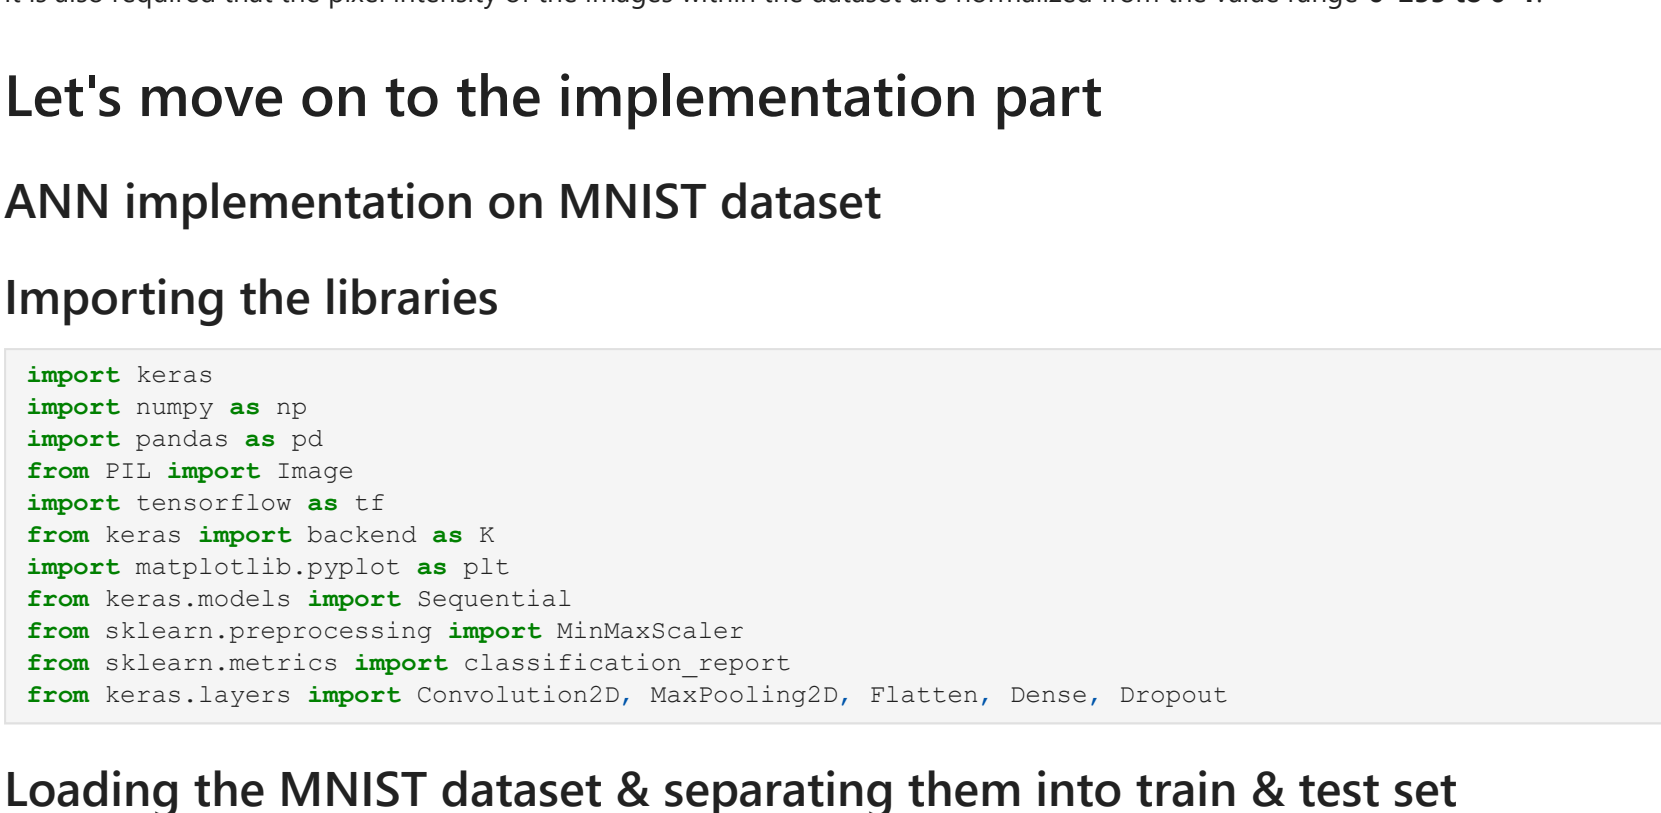
- **Layers**: all the learning occurs in the layers. There are 3 layers
  1. Input
  2. Hidden
  3. Output
- **Feature & Label**: Input data to the network (features) and output from the network (labels)
- **Loss function**: Metric used to estimate the performance of the learning phase
- **Optimizer**: Improve the learning by updating the knowledge in the network

A neural network will take the input data and push them into an ensemble of layers. The network needs to evaluate its performance with a loss function. The loss function gives to the network an idea of the path it needs to take before it masters the knowledge. The network needs to improve its knowledge with the help of an optimizer. Here is an image of various layers and neurons with their connections.

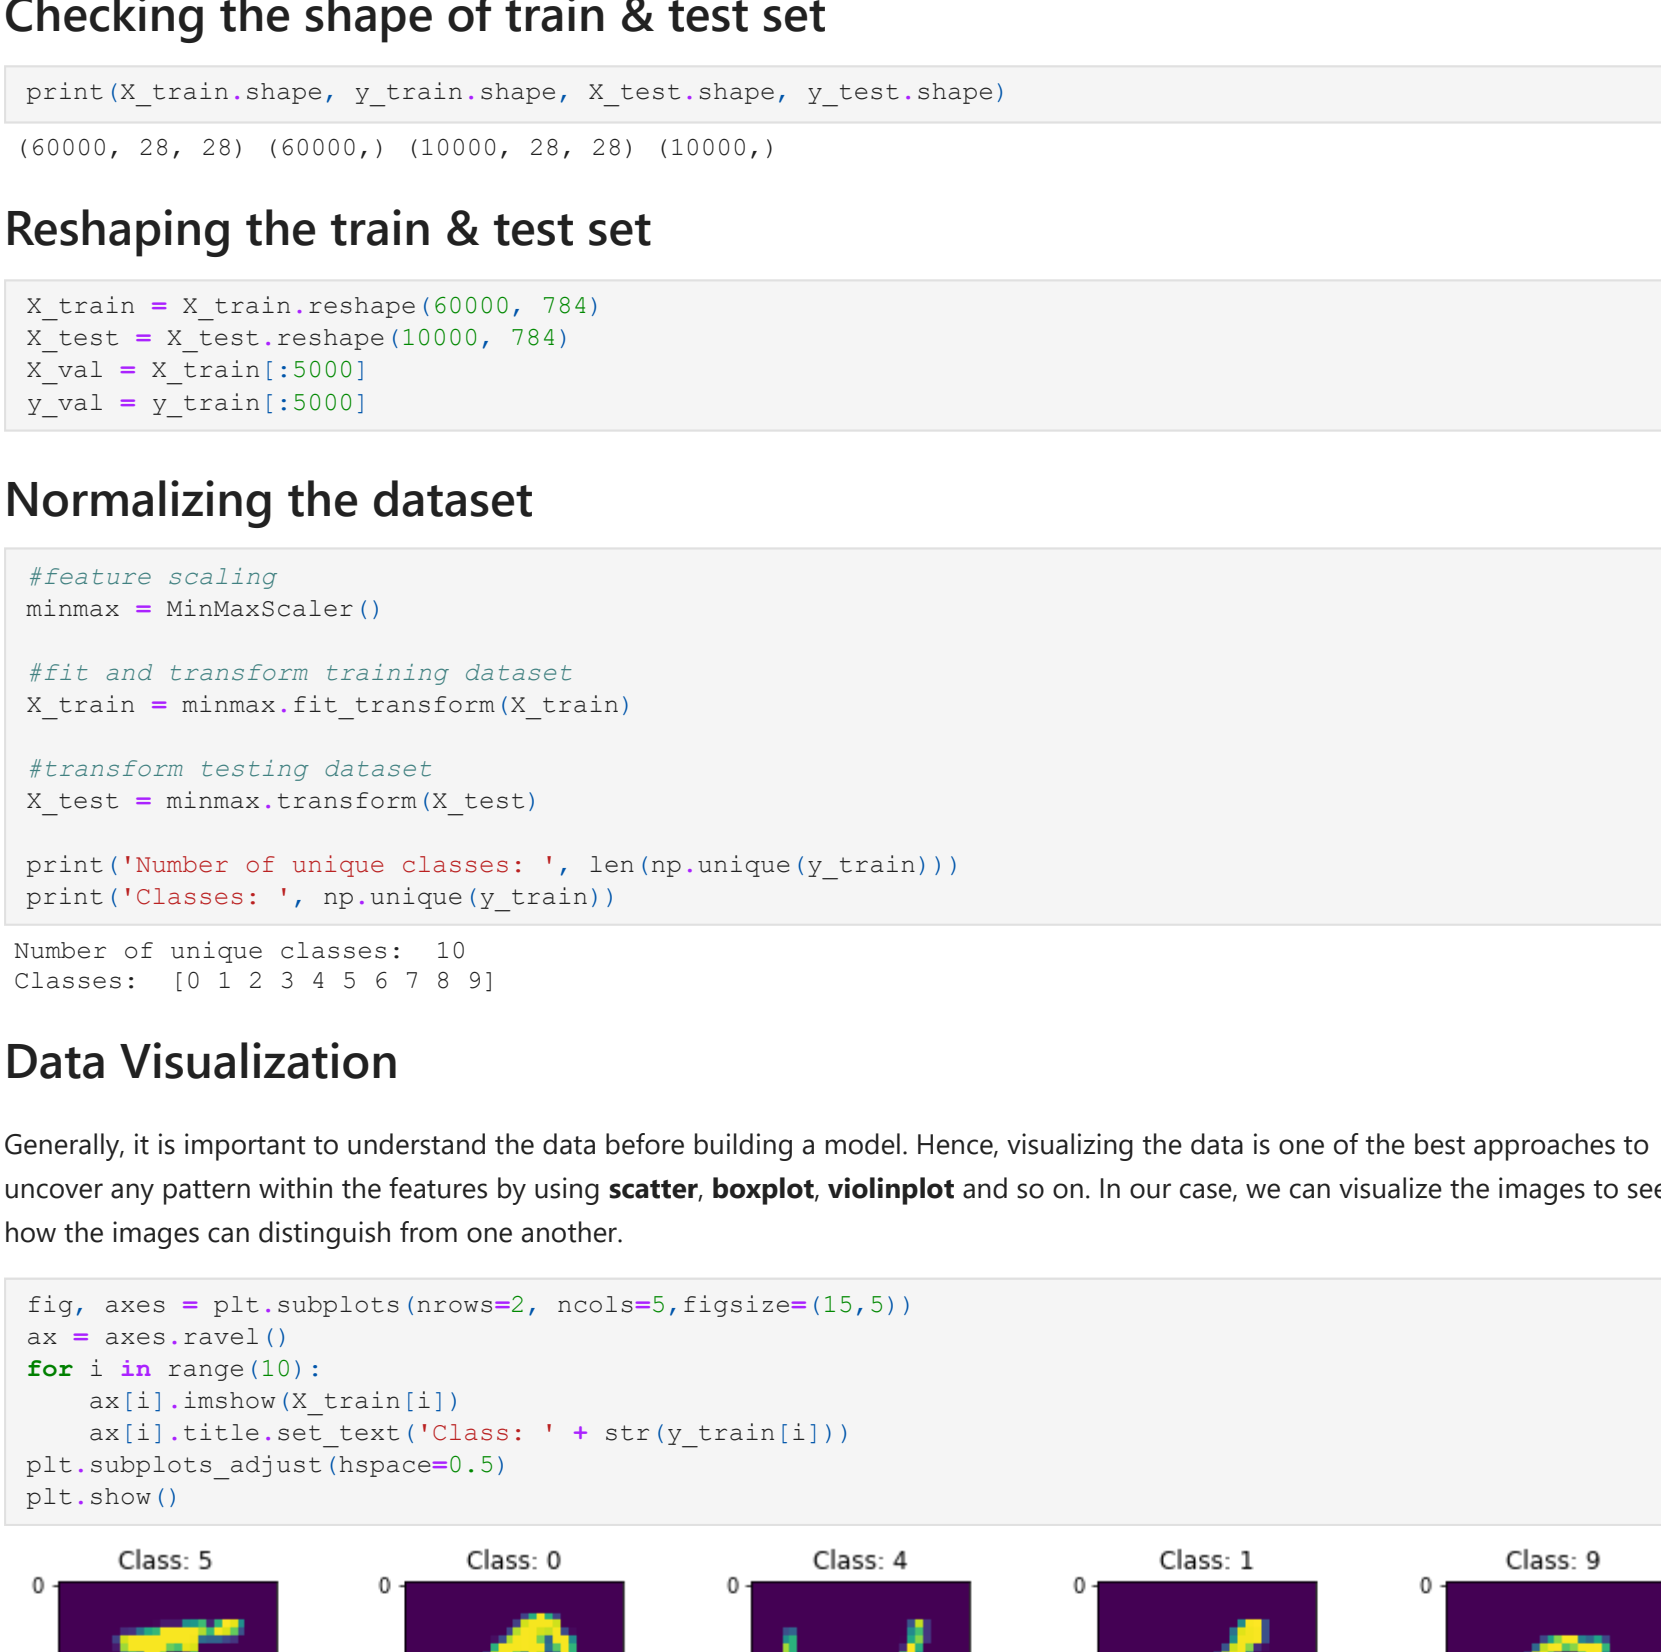


## MNIST

The **MNIST** database (National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems. The database is also widely used for training and testing in the field of machine learning. It was created by "re-mixing" the samples from NIST's original datasets. It consists of the black and white images from NIST which were normalized to fit into a **28x28** pixel bounding box and anti-aliased, which introduced grayscale levels. The **MNIST** database contains **60,000** training images and **10,000** testing images. Half of the training set and half of the test set were taken from NIST's training dataset, while the other half of the training set and the other half of the test set were taken from NIST's testing dataset. The original creators of the database keep a list of some of the methods tested on it. Let's have a look at an example from the dataset. Below is the grayscale image of digit **8** from the dataset.



Below is the frequency of number of image of each digit in the MNIST data.



Now, we load the **MNIST** dataset using the **Keras** library. The **Keras** library has a suite of datasets readily available for use with easy accessibility. We are also required to partition the dataset into **testing**, **validation** & **training**. Here are some quick descriptions of each partition category.

- **Training Dataset**: This is the group of our dataset used to train the neural network directly. Training data refers to the dataset partition exposed to the neural network during training.
- **Validation Dataset**: This group of the dataset is utilized during training to assess the performance of the network at various iterations.
- **Test Dataset**: This partition of the dataset evaluates the performance of our network after the completion of the training phase.

It is also required that the pixel intensity of the images within the dataset are normalized from the value range **0-255** to **0-1**.

## Let's move on to the implementation part

### ANN implementation on MNIST dataset

#### Importing the libraries

```
In [1]: import numpy as np
import pandas as pd
from IPython.display import Image
import tensorflow as tf
from keras import backend as K
import matplotlib.pyplot as plt
from keras.models import Sequential
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import classification_report
from keras.layers import Convolution2D, MaxPooling2D, Flatten, Dense, Dropout
```

#### Loading the MNIST dataset & separating them into train & test set

```
In [2]: mnist = keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

#### Checking the shape of train & test set

```
In [3]: print(x_train.shape, y_train.shape, x_test.shape, y_test.shape)
(60000, 28, 28) (60000, y_train) (10000, 28, 28) (10000, y_test)
```

#### Reshaping the train & test set

```
In [4]: x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_val = x_train[:5000]
y_val = y_train[:5000]
```

#### Normalizing the dataset

```
In [5]: #feature scaling
minmax = MinMaxScaler()

#fit and transform training dataset
x_train = minmax.fit_transform(x_train)

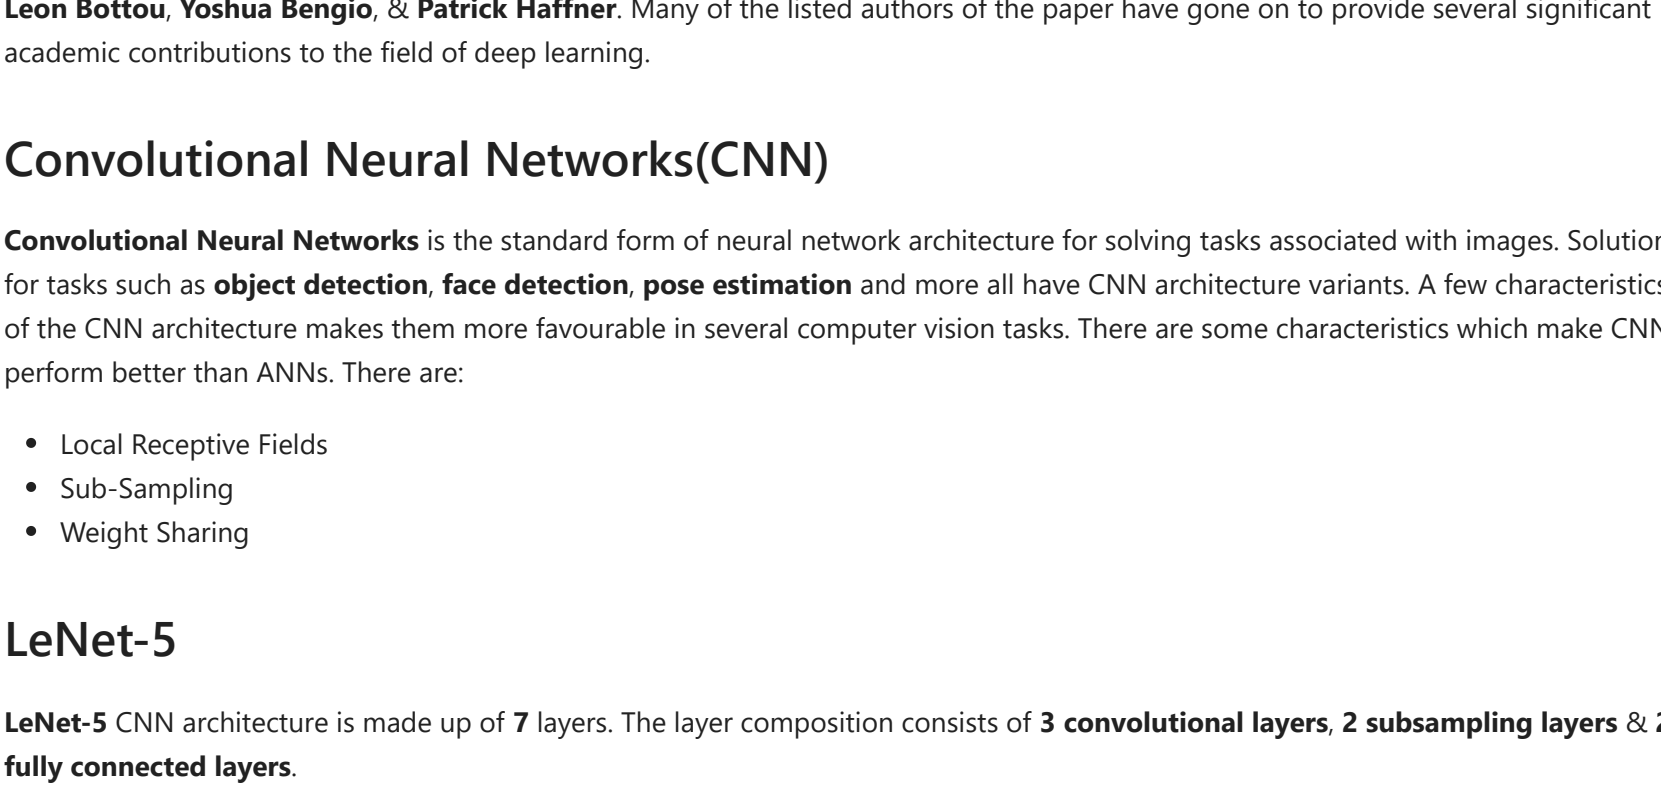
#transform testing dataset
x_test = minmax.transform(x_test)

print('Number of unique classes: ', len(np.unique(y_train)))
print('Classes: ', np.unique(y_train))

Number of unique classes: 10
Classes: [0 1 2 3 4 5 6 7 8 9]
```

#### Data Visualization

Generally, it is important to understand the data before building a model. Hence, visualizing the data is one of the best approaches to uncover any pattern within the features by using **scatter**, **boxplot**, **violinplot** and so on. In our case, we can visualize the images to see how the images can distinguish from one another.



#### Building an ANN model with 2 Dense Layer

```
In [6]: #initializing CNN model
ann_model = Sequential()

#adding 1st hidden layer
ann_model.add(Dense(input_dim = 784, units = 256, kernel_initializer='uniform', activation='relu'))

#adding 2nd hidden layer
ann_model.add(Dense(64, activation='relu'))

#adding output layer
ann_model.add(Dense(units = 10, kernel_initializer='uniform', activation='softmax'))
```

- The **Sequential** class: Sequential groups a linear stack of layers into a **tf.keras.Model**.
- The **add()** method: Adds a layer instance on top of the layer stack.
- The **Dense** class: It adds the hidden layer in our network.
- The **compile()** method: Configures the model for training.
- The **summary()** method: Prints a string summary of the network.

#### Summary of ANN Model

```
In [7]: #model summary
ann_model.summary()

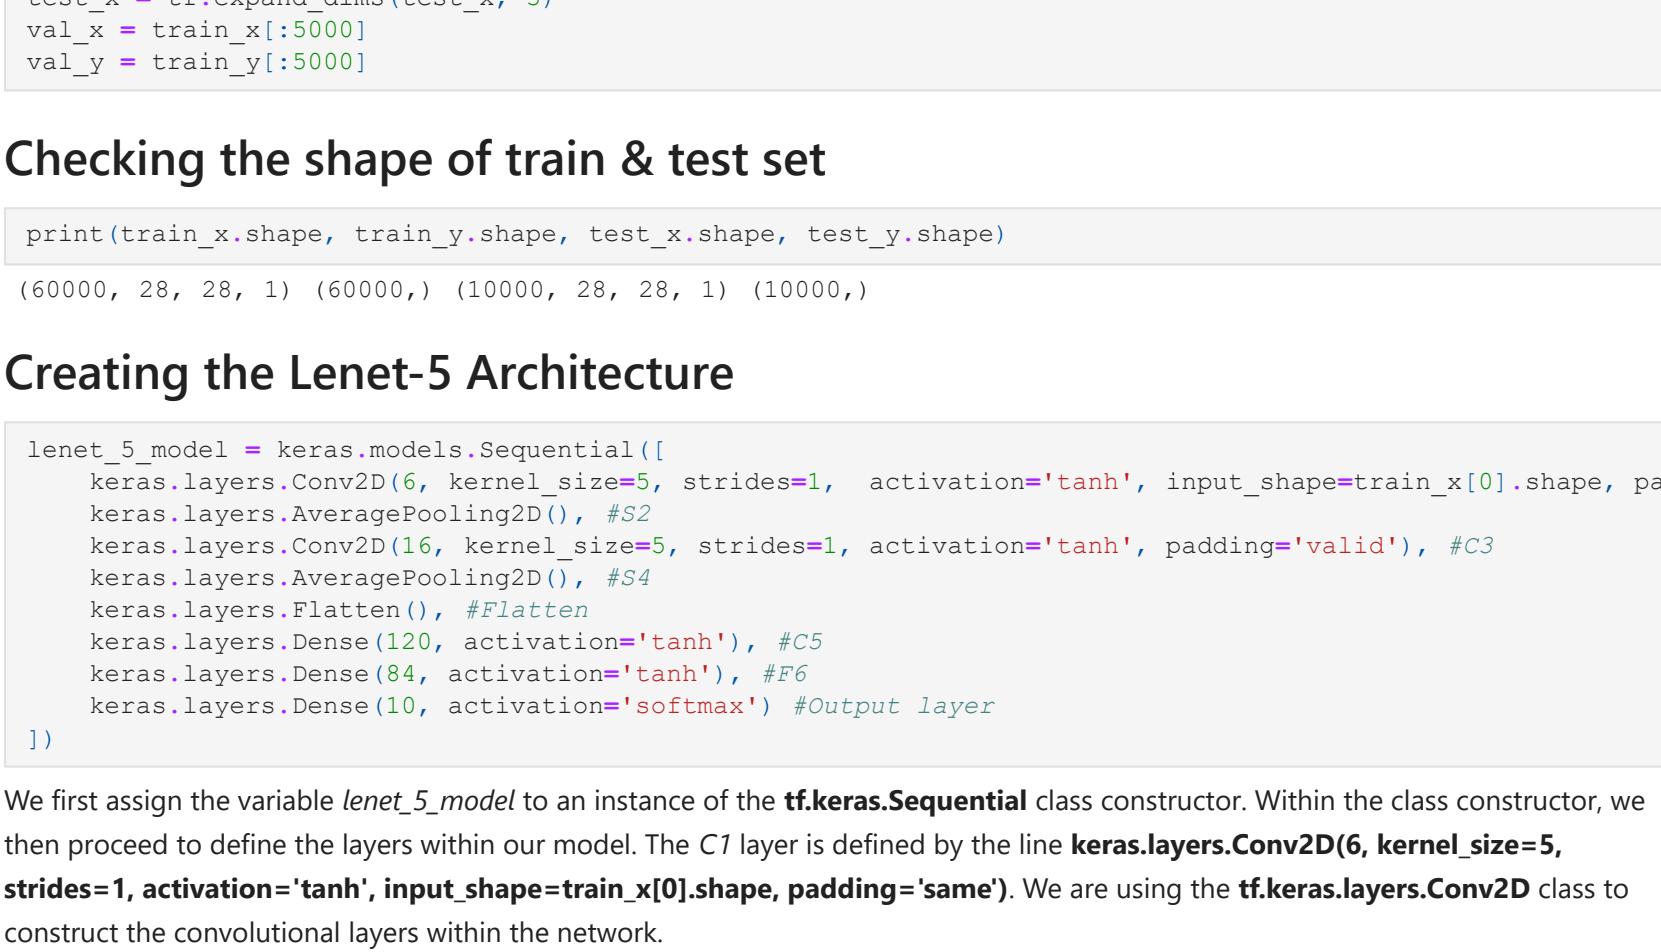
Model: "sequential"
Layer (type) Output Shape Param #
-----
dense_1 (Dense) (None, 256) 200960
dense_2 (Dense) (None, 64) 16448
dense_3 (Dense) (None, 10) 650
Total params: 218,058
Trainable params: 218,058
Non-trainable params: 0
```

#### Compiling the model

```
In [8]: #compile the neural network
ann_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['acc'])
```

#### Training ANN Model

```
In [9]: from tensorflow.keras.callbacks import PlotLossesKeras
ann_model.fit(x_train, y_train, epochs=10, validation_data=(x_val, y_val), callbacks=[PlotLossesKeras()])
```

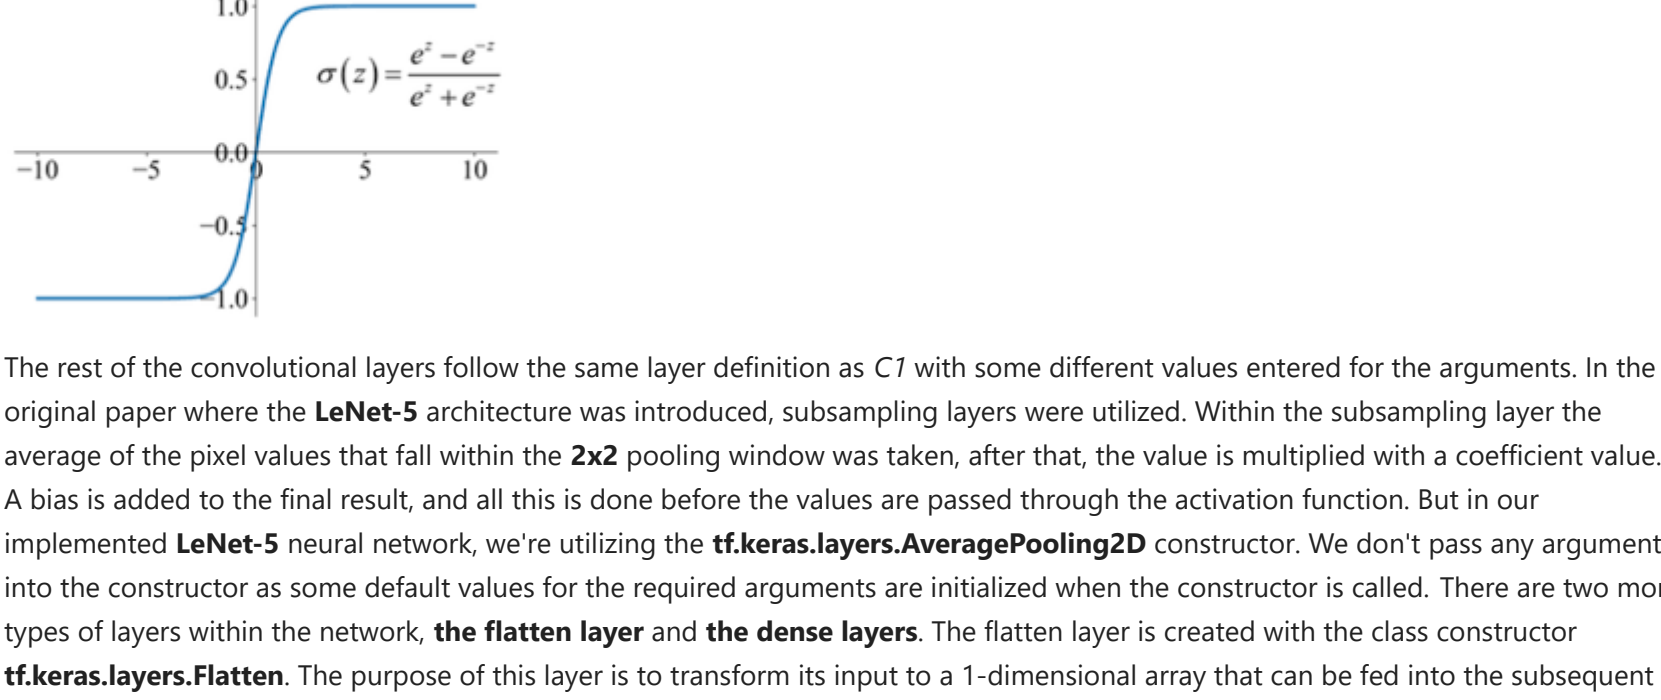


#### Calculating the Accuracy, Precision, Recall & F1-Score of ANN Model

```
In [10]: ann_model.evaluate(x_test, y_test, verbose=0)[1]
```

Out[10]: 0.975000021457621

```
In [11]: y_pred = ann_model.predict(x_test, verbose=1)
y_pred_bool = np.argmax(y_pred, axis=1)
```



**LeNet** was introduced in the research paper *Gradient-Based Learning Applied To Document Recognition* in the year **1998** by **Yann LeCun**, **Leon Bottou**, **Yoshua Bengio**, & **Patrick Haffner**. Many of the listed authors of the paper have gone on to provide several significant academic contributions to the field of deep learning.

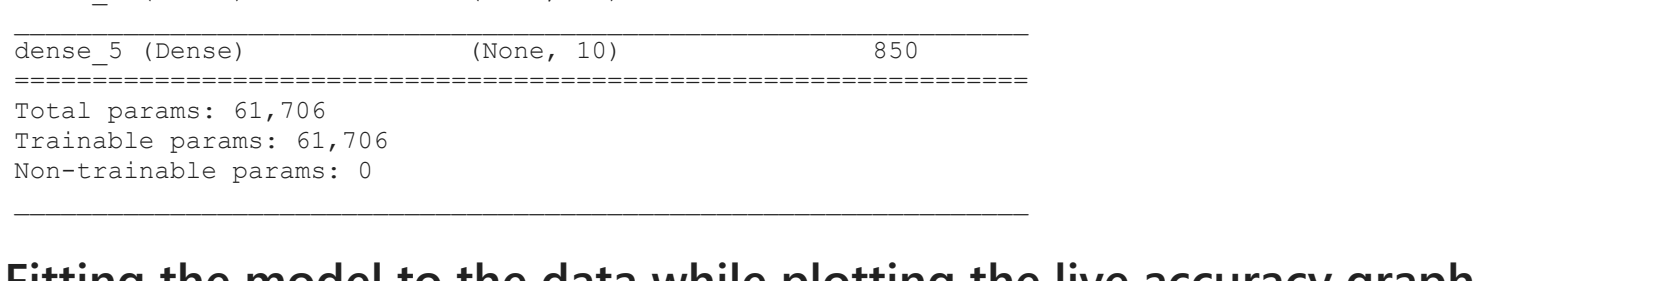
## Convolutional Neural Networks(CNN)

**Convolutional Neural Networks** is the standard form of neural network architecture for solving tasks associated with images. Solutions for tasks such as **object detection**, **face detection**, **pose estimation** and more all have CNN architecture variants. A few characteristics of the CNN architecture makes them more favourable in several computer vision tasks. There are some characteristics which make CNN perform better than ANNs. These are:

- Local Receptive Fields
- Sub-Sampling
- Weight Sharing

## LeNet-5

**LeNet-5** CNN architecture is made up of **7** layers. The layer composition consists of **3 convolutional layers**, **2 subsampling layers** & **2 fully connected layers**.



The diagram above shows a depiction of the **LeNet-5** architecture, as illustrated in the original paper. The proposed model structure of **LeNet-5** has **7** layers, excluding input layer. The details of each layer are as follows:

- **Layer C1** is the first Conv-layer with **6** feature maps with strides of 1. Using a formula given in the appendix, one can calculate the output dimension of this layer **28x28** with **156** trainable parameters. Activation function of this layer is **tanh**.
- **Layer S2** is an average pooling layer. This layer halves average values from the previous Conv layer to the next Conv layer. The Pooling layer is used to reduce the dependence of the model on the location of the features rather than the shape of the features. The pooling layer in LeNet model has a size of **2** and strides of **2**.
- **Layer C3** is the second set of the convolutional layer with **16** feature maps. The output dimension of this layer is **10** with **2,416** parameters. Activation function of this layer is **tanh**.
- **Layer S4** is another average pooling layer with dimension of **2** and stride size of **2**. The next layer is responsible for flattening the output of the previous layer into one dimensional array. The output dimension of this layer is **400** (**5x5x16**).
- **Layer C5** is a dense block (fully connected layer) with **120** connections and **48,120** parameters (**400x120**). Activation function of this layer is **tanh**.
- **Layer F6** is another dense block with **84** parameters and **10,164** parameters (**84x120+84**). Activation function of this layer is **tanh**.
- **Output Layer** has **10** dimension (equals number of classes in the database) with **850** parameters (**10x84+10**). **Activation function of output layer** is **sigmoid**.

The first layer is the input layer — this is generally not considered a layer of the network as nothing is learnt in this layer. The input layer is built to take in **32x32**, and these are the dimensions of images that are passed into the next layer. Those who are familiar with the **MNIST** dataset will be aware that the **MNIST** dataset images have the dimensions **28x28**. To get the **MNIST** images dimension to the meet the requirements of the input layer, the **28x28** images are padded.

The grayscale images used in the research paper had their pixel values normalized from **0** to **255**, to values between **-0.1** & **1.175**. The reason for normalization is to ensure that the batch of images have a mean of **0** and a standard deviation of **1**, the benefits of this is seen in the reduction in the amount of training time. In the image classification with **LeNet-5**, will be normalizing the pixel values of the images to take on values between **0** to **1** so that value will be centered around **0**.

The **LeNet-5** architecture utilizes two significant types of layer construct: convolutional layers and subsampling layers.

- Convolutional layers
- Sub-sampling layers

Within the research paper, convolutional layers are identified with the **Cx**, and subsampling layers are identified with **Sx**, where **x** is the sequential position of the layer within the architecture. **Fx** is used to identify fully connected layers.

The official first layer convolutional layer **C1** produces as output **6** feature maps, and has a kernel size of **5x5**. The kernel/filter is the name given to the window that contains the weight values that are utilized during the convolution of the weight values with the input values. **5x5** is also indicative of the local receptive field size each unit or neuron within the convolutional layer. The dimensions of the six feature maps the first convolution layer produces are **28x28**.

A subsampling layer **S2** follows the **C1** layer. The **S2** layer halves the dimension of the feature maps it receives from the previous layer, this is known commonly as downsampling. The **S2** layer also produces **6** feature maps, each one corresponding to the feature maps as input from the convolutional layer.

```
In [ ]: 
```

## Let's move on to the implementation part

### LeNet-5 TensorFlow Implementation

#### Importing the required libraries

```
In [12]: import tensorflow as tf
from tensorflow import keras
import numpy as np
```

We will be performing the experiment on **MNIST** dataset.

#### Loading the MNIST dataset & separating them into train & test set

```
In [13]: (train_x, train_y), (test_x, test_y) = keras.datasets.mnist.load_data()
train_x = train_x / 255.0
test_x = test_x / 255.0
train_x = tf.expand_dims(train_x, 3)
test_x = tf.expand_dims(test_x, 3)
val_x = train_x[:5000]
val_y = train_y[:5000]
```

#### Checking the shape of train & test set

```
In [14]: print(train_x.shape, train_y.shape, test_x.shape, test_y.shape)
(60000, 28, 28, 1) (60000,) (10000, 28, 28, 1) (10000,)
```

#### Creating the LeNet-5 Architecture

```
In [15]: lenet_5_model = keras.models.Sequential([
    keras.layers.Conv2D(6, kernel_size=5, strides=1, activation='tanh', input_shape=train_x[0].shape, padding='valid'), #C1
    keras.layers.AveragePooling2D(), #S2
    keras.layers.Conv2D(16, kernel_size=5, strides=1, activation='tanh', padding='valid'), #C3
    keras.layers.AveragePooling2D(), #S4
    keras.layers.Flatten(), #F6
    keras.layers.Dense(120, activation='tanh'), #C5
    keras.layers.Dense(84, activation='tanh'), #F6
    keras.layers.Dense(10, activation='softmax') #Output Layer
])
```

We first assign the variable **lenet\_5\_model** to an instance of the **tf.keras.Sequential** class constructor. Within the class constructor, we then proceed to define the layers within the model. The **C1** layer is defined by the line **keras.layers.Conv2D(6, kernel\_size=5, strides=1, activation='tanh', input\_shape=train\_x[0].shape, padding='valid')**. We are using the **tf.keras.layers.Conv2D** class to construct the convolutional layers within the network.

#### Activation Function

A mathematical operation that transforms the result or outputs of neurons into a normalized output. An activation function is a component of a neural network that introduces non-linearity within the network. The inclusion of the activation function enables the neural network to have greater representational power and solve complex functions. In **LeNet-5** architecture they used **tanh** activation function. Below is image of **Tanh** activation.



#### Softmax

An activation function that is utilized to derive the probability distribution of a set of numbers within an input vector. The output of a softmax activation function is a vector in which its set of values represents the probability of an occurrence of a class/event. The values within the vector all add up to 1.

#### Compiling & building the model

```
In [16]: lenet_5_model.compile(optimizer='adam', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
```

**Keras** provides the **compile** method through the model object we have instantiated earlier. The compile function enables the actual building of the model we have implemented behind the scene with some additional characteristics such as the **loss function**, **optimizer**, & **metrics**. To train the network, we utilize a loss function that calculates the difference between the predicted values provided by the network and actual values of the training data. The loss values accompanied by an optimization algorithm (**Adam**) facilitates the number of changes made to the weights within the network. Supporting factors such as **momentum** & **learning rate schedule**, provide the ideal environment to enable the network training to converge, hereby getting the loss values as close to zero as possible. During training, we'll also validate our model after every epoch with the validation dataset partition created earlier.

#### Summary of the model

```
In [17]: lenet_5_model.summary()

Model: "sequential_1"
Layer (type) Output Shape Param #
-----
conv2d_1 (Conv2D) (None, 28, 28, 6) 156
average_pooling2d_1 (AveragePool2D) (None, 14, 14, 6) 0
conv2d_2 (Conv2D) (None, 10, 10, 16) 2416
average_pooling2d_2 (AveragePool2D) (None, 5, 5, 16) 0
flatten_1 (Flatten) (None, 400) 0
dense_1 (Dense) (None, 120) 48120
dense_2 (Dense) (None, 84) 8154
dense_3 (Dense) (None, 10) 100
Total params: 61,706
Trainable params: 61,706
Non-trainable params: 0
```

#### Fitting the model to the data while plotting the live accuracy graph

```
In [21]: from tensorflow.keras.callbacks import PlotLossesKeras
ann_model.fit(train_x, train_y, epochs=10, validation_data=(val_x, val_y), callbacks=[PlotLossesKeras()])
```



#### Calculating the Accuracy, Precision, Recall & F1-Score of Lenet-5

```
In [22]: lenet_5_model.evaluate(test_x, test_y)[1]
```

Out[22]: 0.987800002980835

```
In [23]: y_pred = lenet_5_model.predict(test_x, verbose=1)
y_pred_bool = np.argmax(y_pred, axis=1)
```



#### Analysis of Lenet architecture

**LeNet-1** was trained on LeCun's USPS datasets, whereas **LeNet-4** & **LeNet-5** were trained on the MNIST dataset. All LeNet variants were commonly trained using **Maximum Likelihood Estimation (MLE)** and **Mean Squared Error (MSE)** loss. The input images for LeNet were normalised such that their values remain within the range **[+0.1, 1.175]**, which made the mean **0** and the variance roughly **1**. It was argued that normalisation accelerated the network's learning. While the LeNet CNNs marked a breakthrough for their time, there were several limitations:

1. The networks were **small** and therefore had limited applications where they worked correctly.
2. The networks only worked with **single-channel (black and white images)**, which also limited its applications.
3. Most modern adaptations of this model implement a **max-pooling operation** as opposed to an **average pooling** operation as it allows for more rapid convergence due to the larger gradients returned during back-propagation.
4. Most modern adaptations implement a **Relu activation** as opposed to **tanh** and **sigmoid** as **Relu** usually leads to higher classification accuracies.

#### Summary of LeNet

However, the following represent some key takeaways about the LeNet networks:

1. Operated with single-channel images, either **(28x28)** or **(32x32)** in size.
2. Used **tanh** activation functions as opposed to **Relu**.
3. Used **average pooling** as opposed to **max pooling** for subsampling.
4. Used **(2x2)** kernels for average pooling subsampling.
5. Used **(5x5)** kernels for convolutions.
6. Trained by optimising for **Mean Squared Error (MSE) loss**.
7. Included **trainable coefficient** and **bias** for some subsampling operations.
8. Classified images into one of ten classes.

## General & Architectural comparison between CNNs(LeNet-5) & ANNs

Artificial Neural Networks (ANN)	Convolutional Neural Networks (CNN)
<b>ANN</b> consist of dense layers where all the input data is sent to each neuron on the first layer which then performs a dot product of the input data and the neuron's weights to produce a single number as output	In <b>CNN</b> , the weights (in the convolutional layers) are a small matrix (often 3x3) which is dot-producted with each pixel to produce a new pixel thus acting as image filters
<b>ANN</b> is sometimes referred to as a Feed-Forward Neural Network because inputs are processed only in a forward facing direction	<b>CNN</b> doesn't process data in a forward-facing way, but rather refers to the input data multiple times when creating map
With <b>ANN</b> , concrete data points must be provided. For example, in a model where we are trying to distinguish between dogs and cats, the width of the noses and length of the ears must be explicitly provided as data points	When using <b>CNN</b> , these spatial features are extracted from image input. Instead of <b>CNN</b> ideal when thousands of features need to be extracted. It's more of having to measure each individual feature. <b>CNN</b> gathers these features from its own
<b>ANN</b> loses the spatial relationship between the pixels of the images as we have to flatten all the pixels	<b>CNN</b> keeps the spatial information in the image hence perform better than <b>ANNs</b>
In <b>ANN</b> , number of trainable parameters are way more as each neuron is fully connected to all the neurons of next layer	In <b>CNN</b> , number of parameters are relatively smaller as the weights get shared when we use kernel for <b>Convolution</b>
<b>ANN</b> are not translational invariant	<b>CNN</b> are translational invariant due to application of Pooling operation while training
<b>ANNs</b> are not very efficient for images inspite of being expensive	<b>CNNs</b> are very efficient for images & very less expensive compared to <b>ANN</b>
Number of parameters of my ANN architecture are around 218K	Number of parameters of my <b>LeNet-5</b> architecture are around 61K
Overall time taken to train the ANN on my system for 10 epochs was around 110s	Overall time taken to train <b>LeNet-5</b> on my system for 10 epochs was around 80s

&lt;



